

Input/Output for QCD

Version 2.0

SciDAC Software Coordinating Committee

May 4, 2006

1 Introduction

This document describes the Input/Output Applications Programmer Interface developed under the auspices of the U.S. Department of Energy Scientific Discovery through Advanced Computing (SciDAC) program.

Although the QIO I/O system was developed to support the data parallel lattice-QCD API called QDP/C and QDP++, it is designed to function independently of QDP, requiring only the lower level QMP message passing package. Two data models are treated: lattice fields, consisting of data of the same format residing on each site of a hypercubic lattice, and global data, constant across all lattice sites. Lattice field data is distributed among multiple nodes.

The file format consists of a series of logical records. Each record contains user-controlled metadata and binary data. An arbitrary combination of logical records is permitted. The physical file format is based on a custom SciDAC LIME standard, (Lattice QCD Interchange Message Encapsulation), which views the file as a series of LIME messages, each, in turn, consisting of a series of LIME records. Details of the physical format are hidden from the user. The LIME package is included with QIO.

2 Overview of Binary File Format

2.1 Introduction

The binary file format has been designed with flexibility in mind. For archiving purposes, the allowable logical records, metadata, and binary content may be further restricted. Here we described the unrestricted format.

Three classes of file volumes are supported: single-file volumes, partition-file volumes and multiple-file volumes. Single files are read and written through a single master node; partition files are read and written through a set of designated I/O nodes; and multiple files, through each node. With partition-file and multiple-file formats the binary data is split into separate files, one for each node that participates in I/O. With the single-file format, data is contained in

a single file. Single-processor utilities are provided for converting between the partition-file and single-file formats.

2.1.1 Single file format

Single binary files are composed of a series of one or more logical application records. A single logical record encodes a single lattice field, an array of lattice fields of the same data type, or an array of global data. Physics metadata, managed at the convenience of the applications programmer, is associated with the file itself and separately with each logical record as well. The applications programmer views the file as follows:

- File physics metadata
- Record 1 physics metadata and binary data
- Record 2 physics metadata and binary data
- etc.

For example, a file might record a series of staggered fermion eigenvectors for a gauge field configuration. Each record would then map to a single field for a color vector. The file metadata might include information about the parent gauge field configuration and the record metadata might encode the eigenvalue and an index for the eigenvector.

For another example, the gauge field configuration in four dimensions is represented as an array of four color matrix fields. The configuration is conventionally written so that the four color matrices associated with each site appear together. A file containing a single gauge field configuration would then consist of a single logical record containing the array of four color matrices.

Additional metadata is automatically managed by QIO (without requiring intervention by the applications programmer) to facilitate the implementation and to check data integrity. Thus the file actually begins with private QIO metadata and physics metadata and each logical application record consists of four LIME records. Within QIO the file is viewed as a series of LIME records as follows:

- Private file QIO metadata
- User file physics metadata
- Record 1 private QIO metadata
- Record 1 user physics metadata
- Record 1 binary data
- Record 1 private checksum
- Record 2 private QIO metadata

- Record 2 user physics metadata
- Record 2 binary data
- Record 2 private checksum
- etc.

The site order of the binary data is lexicographic according to the site coordinate r_i with the first coordinate r_0 varying most rapidly. The byte layout of the site data is determined by user-supplied “factory” functions. Physical byte ordering of IEEE numeric data (integers and floating point) is big-endian, regardless of the architecture of the processor that creates the file. To achieve this result, numeric site data within a given logical record must consist of a series of words of the same size, such as an $SU(3)$ matrix represented entirely by single precision words (or entirely double precision). Mixed precision structures are excluded, but 32-bit integers and floats may coexist.

2.2 Partition file format

With cluster computers consisting of many hundreds of processors it may prove impractical to provide NFS mounts from each processor to a common file system. Instead processors can be grouped into I/O sets, each with a single I/O node and disk. Input files are fragmented and staged to disks attached to these nodes and output files are reassembled from them. Single-processor utilities are provided for file disassembly and assembly. It is intended that the installation and implementation hide these details from the user, so the user’s view of the file is the same as with single file format. Ideally the local software environment is designed with portability in mind, so that user code calling for file input or output will get the same result in the end without any change in the code, whether or not the intermediate volume format happens to be single file at one installation and partitioned at another.

Each partition I/O node processes data only for sites stored on that partition. The data for the binary field is divided accordingly, so each partition file holds field data only for the nodes on its partition. For ease in conversion to and from single file format, the site data for a file is always arranged in a standard “lexicographic” order according to the lattice coordinate. The lexicographic rank identifies the site. A binary site list is placed at the beginning of every file to identify its contents. The master I/O node, which must also be a partition I/O node, handles all of the global data in the file, including file and record metadata and any global binary data. Its file format is identical to the single file format, except for the addition of a sitelist record. The other partition files contain only site lists and the binary data for the relevant partition.

All component files are given unique names, constructed by attaching the file extension `.volnnnn`, where `nnnn` is the number of the node that reads the file (with leading zeros). The file is known to the user by its unextended name.

The principal file read by the master node contains most of the metadata.

- Private file QIO metadata
- User file physics metadata
- Binary index of sites
- Record 1 private QIO metadata
- Record 1 user physics metadata
- Record 1 binary data
- Record 1 private checksum
- Record 2 private QIO metadata
- Record 2 user physics metadata
- Record 2 binary data
- Record 2 private checksum
- etc.

The secondary files contain these records:

- Binary index of sites
- Record 1 binary data
- Record 2 binary data
- etc.

The site index in each case is a table of contents, that is, a list of the lexicographic ranks of all sites contained in the file in the order of appearance.

2.3 Multifile format

The API provides for rapid temporary writing of data to scratch disks and reading from scratch disks. In this case it is assumed that the files are not intended for longer term storage. The file formats are identical to the partition file formats with one exception: the site order is internal storage order, rather than lexicographic order. This choice is made to reduce cache-misses during I/O.

3 Metadata Standard and Manipulation

The QIO implementation uses an XML encoding for its private file and record metadata. It is hidden above the QIO API. The data is available to the user through a C structure with accessor functions for retrieving and setting values.

Since QIO processes the user file and record metadata blindly as a character string, QIO places no restrictions on the format of the user metadata.

4 QIO API

This section describes the QIO interface.

The QIO system provides for binary file operation for writing and reading lattice fields and global data. Lattice fields consist of any data type homogeneous over lattice sites or an array of such data types. Global data consists of an array of data types or of strings. The storage of lattice data on the nodes is described in a `QIO_Layout` structure, and the information required for presenting field data in the correct byte order is encapsulated in “factory” functions.

4.1 The layout structure

The structure is defined as follows:

```
typedef struct {
    /* Data distribution */
    int (*node_number)(const int coords[]);
    int (*node_index)(const int coords[]);
    int (*num_sites)(int node);
    void (*get_coords)(int coords[], int node, int index);
    int *latsize;
    int latdim;
    size_t volume;
    size_t sites_on_node;
    int this_node;
    int number_of_nodes;
} QIO_Layout;
```

The data distribution (layout) structure has nine members. The `node_number` member is an implementer-supplied function returning the number of the node that has the specified lattice coordinate. The `node_index` member returns the storage order index for the site on its node. The `get_coords` member maps the node number and index values to lattice coordinates. The `num_sites` member returns the number of sites on the specified node. The next two members specify the lattice coordinate extent and spacetime dimensionality. The seventh member specifies the full spacetime volume. The eighth, the number of sites on the current node, the ninth, the number of the present node, and the ninth, the total number of nodes.

Here is an illustration of how the layout structure is loaded from the data in our implementation of the QDP/C API prior to a `QIO_open_read` or `QIO_open_write` call:

```
QIO_Layout layout;

layout.node_number = QDP_node_number;
layout.node_index  = QDP_index;
layout.get_coords  = QDP_get_coords;
```

```

layout.num_sites = QDP_num_sites;
layout.latdim = QDP_ndim();
layout.latsize = (int *)malloc(layout->latdim*sizeof(int));
QDP_latsize(layout.latsize);
layout.volume = QDP_volume();
layout.sites_on_node = QDP_sites_on_node;
layout.this_node = QDP_this_node;
layout.number_of_nodes = QDP_numnodes();

```

4.2 Private Record Metadata

Field data is described by a private QIO record metadata structure. On output the application must create and populate the structure. On input, the structure is populated from the file.

The private QIO record metadata is used for consistency checking and for providing the user a standard tool for recording and discovering the data type being stored. Semantically, it serves the same purpose as a BinX record. It carries enough information to completely define the binary record format. The record metadata is held in an opaque `QIO_RecordInfo` structure. Elements are accessed and manipulated through the following functions.

Create and populate the private record metadata structure Before writing a record the calling program must create the private record metadata structure. Before reading a record, the calling program must allocate space for the private record metadata structure using the same calling procedure.

Prototype	<code>QIO_RecordInfo *QIO_create_record_info(int globaltype, char *datatype, char *precision, int colors, int spins, int typesize, int datacount);</code>
Example	<code>rec_info = QIO_create_record_info(QIO_FIELD, "QDP_F_Real", "F", 0, 0, size, 1);</code>
Example	<code>rec_info = QIO_create_record_info(0, "", "", 0, 0, 0, 0);</code>

The first example is appropriate for output. The second, for input.

The `globaltype` parameter distinguishes between a record containing a lattice field and a record containing a lattice constant array.

`QIO_FIELD`, `QIO_GLOBAL`

for a field and global constant record type, respectively.

The `datatype` string is not interpreted by QIO. It allows the applications programmer a standard way to identify the data type. For that purpose the name should be unique. For QDP/C we use the datatype name of the QDP field. For global data we use the name of one of the QLA datatypes.

The `precision` string is one of these:

F	single
D	double
S	random number generator state consisting of 32-bit floats and ints
I	integer (currently only 32-bit is supported)

This string is interpreted by the host file conversion utility.

The **colors** and **spins** arguments give the working value for these quantities, if they apply to the datatype. Otherwise, they should be zero. They are not interpreted by QIO.

The **typesize** specifies the number of bytes per site item and the **datacount** specifies the number of such items per site. The product is the total number of bytes per site. For example, for a single precision SU(3) gauge field with four color matrices per site, the typesize is 72 and the datacount is 4.

It is not an error to create a structure with zeros for integer values and null string pointers. Those data items are tagged as “missing”. However, **QIO_write** and **QIO_read** return an error condition, if the total byte count per site is inconsistent with the values in this structure.

Destroy the private record metadata structure

Prototype	<code>void QIO_destroy_record_info(QIO_RecordInfo *record_info);</code>
Example	<code>QIO_destroy_record_info(rec_info);</code>

Compare two private record metadata structures To allow for verification that a record being read matches what is expected, the calling program may create the record information structure that it expects and compare it with the structure that was read from the file.

Prototype	<code>int QIO_compare_record_info(QIO_RecordInfo *found, QIO_RecordInfo *expect);</code>
Example	<code>int ok = QIO_compare_record_info(rec_info, cmp_info);</code>

The arguments are *not* symmetric. Only those fields that are non-empty in the **expect** structure are compared with fields in the **found** structure.

Extract values from the file reader structure The following accessors perform self-evident functions:

Prototype	<code>int QIO_get_reader_latdim(QIO_Reader *in); int *QIO_get_reader_latsize(QIO_Reader *in); uint32_t QIO_get_reader_last_checksuma(QIO_Reader *in); uint32_t QIO_get_reader_last_checksumb(QIO_Reader *in);</code>
-----------	--

Extract values from the file writer structure The following accessors perform self-evident functions:

Prototype	<pre>uint32_t QIO_get_writer_last_checksuma(QIO_Writer *out); uint32_t QIO_get_writer_last_checksumb(QIO_Writer *out);</pre>
-----------	--

Extract values from the private record metadata structure The following accessors perform self-evident functions:

Prototype	<pre>int QIO_get_globaldata(QIO_RecordInfo *record_info); char *QIO_get_datatype(QIO_RecordInfo *record_info); char *QIO_get_precision(QIO_RecordInfo *record_info); int QIO_get_colors(QIO_RecordInfo *record_info); int QIO_get_spins(QIO_RecordInfo *record_info); int QIO_get_typesize(QIO_RecordInfo *record_info); int QIO_get_datacount(QIO_RecordInfo *record_info); char *QIO_get_record_date(QIO_RecordInfo *record_info);</pre>
-----------	--

4.3 Opening and closing binary files

The file opening procedures differ, depending on whether the file is opened for reading or writing.

Open a file for writing

Prototype	<pre>QIO_Writer *QIO_open_write(QIO_String *xml_file, char *filename, int volfmt, QIO_Layout *layout, QIO_Oflag *oflag);</pre>
Purpose	Opens a named file for writing and writes the file metadata.
Example	<pre>QIO_Writer *outfile; QIO_Layout layout; outfile = QIO_open_write(xml_file_out, filename, QIO_SINGLEFILE, &layout, QIO_SERIAL QIO_TRUNC);</pre>

The `QIO_Writer *` return value points to the file handle used in subsequent references to the file. The first argument is the user file XML. To create the `QIO_String` structure, starting from a plain character array, use the command

Prototype	<code>void QIO_string_set(QIO_String *qs, const char *const string)</code>
Example	<pre>QIO_String *xml_file = QIO_string_create(); QIO_string_set(xml_file, xmlstring);</pre>

The next-to-last argument is the layout structure. It is assumed that the user has prepared it as described above.

The `QIO_Oflag` structure is defined as follows:


```
typedef struct {
    int serpar;    /* Placeholder for specifying serial or
                   parallel access */
    int mode;      /* QIO_TRUNC or QIO_APPEND */
} QIO_Oflag;
```

It contains parameters that control whether the file is to be truncated or data is to be appended and has a placeholder for future use for specifying whether the file is to be written in parallel or serially. At present only serial writing is supported. The structure is initialized as in the following example:

```
QIO_Oflag oflag;
oflag.serpar = QIO_SERIAL;
oflag.mode   = QIO_TRUNC;
```

These are the default values used when the `&oflag` parameter is passed as a null pointer.

Caution: If a file is opened for appending, QIO presently does not verify that the fields being appended conform to the lattice dimensions and layout of the fields already present.

Open a file for reading

Prototype	<code>QIO_Reader *QIO_open_read(QIO_String *xml_file, char *filename, QIO_Layout *layout, QIO_Iflag *iflag);</code>
Purpose	Opens a named file for reading and reads the file metadata.
Example	<code>QIO_Reader *infile; QIO_Layout layout; infile = QIO_open_read(xml_file_in, filename, &layout, QIO_SERIAL);</code>

The `QDP_Reader` return value is the file handle used in subsequent references to the file. A null return value signals an error. It is assumed the user has created the file metadata structure with address `xml_file`, so it can be read from the head of the file and inserted. Space for the string within the structure is reallocated to a sufficient size by QIO. The other arguments have the same meaning as with `QIO_open_write`. The volume format is auto-detected so is not specified by the calling program. It is assumed that the user has prepared the `layout` argument as described above.

The `QIO_Iflag` structure is defined as follows:

```
typedef struct {
    int serpar;    /* Placeholder for specifying serial or
                   parallel access */
    int volfmt;    /* QIO_UNKNOWN, QIO_SINGLEFILE, QIO_PARTFILE,
                   QIO_MULTIFILE */
} QIO_Iflag;
```

A file is usually opened with automatic detection of the file format. However, confusion arises when the file appears in both formats in the same directory. In that case the `volfmt` member is needed to specify a preference. Otherwise, the parameter can be safely passed as `QIO_UNKNOWN` or `QIO_SINGLEFILE`, regardless of the file format, and the format will be set according to the existing file. The structure also has a placeholder for future use for specifying whether the file is to be read in parallel or serially. At present only serial reading is supported. The structure is initialized as in the following example:

```
QIO_Iflag iflag;
iflag.serpar = QIO_SERIAL;
iflag.mode   = QIO_UNKNOWN;
```

These are the default values used when the `&iflag` parameter is passed as a null pointer.

In normal operation the user specifies the lattice dimension in the `QIO_Layout` structure, and an error condition occurs, if the dimensions in the file do not match the dimensions in the layout structure. Provision is made to operate in discovery mode. If the layout `latdim` member is zero when `QIO_open_read` is called, no checking takes place and the lattice dimensions are taken from the file and kept with the `QIO_Reader` structure. The user's `QIO_layout` structure is not altered by QIO. Instead, it works with an updated internal copy of that structure, kept in the opaque `QIO_Reader`. Two accessor functions are provided for extracting the dimensions from the reader:

Get the number of spacetime dimensions

Prototype	<code>int QIO_get_reader_latdim(QIO_Reader *in);</code>
Purpose	Returns the number of spacetime dimensions.
Example	<code>int latdim = QIO_get_reader_latdim(qio_in);</code>

Get the lattice size in each direction

Prototype	<code>int *QIO_get_reader_latsize(QIO_Reader *in);</code>
Purpose	Returns a pointer to an integer array of sizes for each dimension.
Example	<code>int *latsize = QIO_get_reader_latsize(qio_in);</code>

Allocation of the array is controlled by QIO. The array storage is released by the `QIO_close_read` call.

Close an output file

Prototype	<code>int QIO_close_write(QIO_Writer *out);</code>
Example	<code>QIO_close_write(outfile);</code>

Close an input file

Prototype	<code>int QIO_close_read(QIO_Reader *in);</code>
Example	<code>QIO_close_read(infile);</code>

In both cases the integer return value is 0 for success and 1 for failure.

4.4 Writing and reading fields, arrays of fields, or arrays of global data

Prototype	<code>int QIO_write(QIO_Writer *out, QIO_RecordInfo *record_info,QIO_String *xml_record, void (*get)(char *buf, size_t index, size_t count, void *arg), int datum_size, int word_size, void *arg);</code>
Example	<code>QIO_RecordInfo *rec_info; rec_info = QIO_create_record_info(QDP_FIELD,"QDP_F_Real","F", 0,QLA_Ns,size,1); QIO_write(outfile, rec_info, xml_record, QDP_F_get_R, sizeof(QLA_Real), sizeof(QLA_Real), (void *)field);</code>

The integer return value is 0 for success and 1 for failure. It is assumed the user has prepared the record metadata and the field data in advance.

The input arguments are as follows:

<code>out</code>	The <code>QIO_Writer</code> handle returned by <code>QIO_open_write</code> .
<code>record_info</code>	The private metadata for the record (see below).
<code>xml_record</code>	The user-constructed metadata for the record.
<code>get</code>	Factory function (see below).
<code>datum_size</code>	The total number of bytes required to serialize the datum.
<code>word_size</code>	The number of bytes in a datum word.
<code>arg</code>	Pass-through parameters for the factory function.

The second argument, the `record_info` structure, contains information about the data format, as described in Sec. 4.2. It must be created by the caller in all cases. For output, the caller must set its values. For input, the values are returned from the file.

The fourth argument is a factory function that, in this example, is invoked by QIO like this:

```
QDP_F_get_R(buf, index, count, field);
```

It is expected to fill the QIO-supplied buffer `buf` with a byte-serialized copy of the field datum at site index `index`. The parameter `count` specifies the array length of the field datum at that site. The datum size parameter `datum_size` gives the total number of bytes to be delivered as the product of the count parameter and the byte length of the array element on that site.

It is up to the applications programmer to insure that the data base-type (int, float, double) word order produced by the factory function follows the

SciDAC convention for the specified datatype. However byte ordering within a word (big endian or little endian) processed by the factory functions should be in the native order of the architecture. Any byte rearrangement needed to convert to and from standard file endianness is the responsibility of QIO. To this end the user must specify the base-type word length of the data in bytes through the parameter `word_size`. All numeric SciDAC data types are homogeneous in word size, so a single parameter suffices.

For example for an array of four single precision color vector fields, each consisting of three complex numbers, there are $4 \times 3 \times 2 = 24$ real values per site, each of them single-precision floating point numbers. The word size for the IEEE float datatype is 4 (bytes). The factory function must produce the standard word order: real part of the first color component of the first color vector, followed by the imaginary part of the same component, followed by the real and then imaginary parts of the second color component of the first color vector, etc. The count is 4 (color vectors), and the datum size is $4 \times 24 = 96$ (total bytes per call). [The type size of $3 \times 2 \times 4 = 12$ (bytes) and the count of 4 (array elements) were specified when creating the `QIO_record_info` structure.]

The same factory function signature is used for global and field data, even though for global data the site `index` parameter has no meaning. The applications programmer would doubtless provide different functions for the two cases. For field data, QIO calls the factory function once per lattice site. For global data, QIO calls only once and expects to take all the data in that call. It is the responsibility of the applications programmer to provide the appropriate factory function for each case.

Since the open operation has already registered a `node_number` function, QIO knows to ask only for a site on the present node. The factory function is not required to fetch data from a different node.

The seventh argument of `QIO_write` is passed through as the fourth argument of the `get` function. It can be used to identify the field from which the data is required. In this way only one factory function is needed for each QDP and QLA datatype.

Read a field, array of fields, or array of global data

Prototype	<pre>int QIO_read(QIO_Reader *in, QIO_RecordInfo *record_info, QIO_String *xml_record, void (*put)(char *buf, int coords[], void *arg), int datum_size, void *arg);</pre>
Example	<pre>QIO_read(infile, rec_info, xml_record, QDP_F_put_r, sizeof(QLA_Real), (void *)field);</pre>

The integer return value is 0 for success and 1 for failure. It is assumed the user has prepared the record metadata and the field data in advance. This operation is the inverse of the write operation described. The `put` factory function does the reverse of the `get` function.

Read only the record metadata This utility makes it possible to examine only the header of the record in order to decide whether to continue reading. The state of the file is remembered, so a subsequent call to `QIO_read` reads the full record as though this call had not been made.

Prototype	<code>int QIO_read_record_info(QIO_Reader *in, QIO_RecordInfo *record_info, QIO_String *xml_record);</code>
Example	<code>QIO_read_record_info(infile, rec_info, xml_record);</code>

Skip to the next record

Prototype	<code>int QIO_next_record(QIO_Reader *in);</code>
Example	<code>QIO_next_record(infile);</code>

Set and determine the verbosity level A user can control the verbosity of QIO. Choices in increasing degree of chatter are

`QIO_VERB_OFF`
`QIO_VERB_LOW`
`QIO_VERB_MED`
`QIO_VERB_REG`
`QIO_VERB_DEBUG`

Prototype	<code>int QIO_verbosity(int level);</code>
Example	<code>oldlevel = QIO_verbosity(QIO_OFF);</code>

A user can also inquire about the current verbosity level with the following function.

Prototype	<code>level = QIO_verbosity();</code>
-----------	---------------------------------------

4.5 File format conversion

The API provides subroutines for converting between single file and partition file format. Since the partition file format depends on which nodes are I/O nodes and it depends on the data layout as it appears on the compute nodes, the complete code for carrying out file conversion requires an implementation suited to the locale.

The file conversion utilities require information about the data layout on the compute nodes. This information is provided by the `QIO_Layout` structure as described above. Furthermore, it requires information about the file system and the identity of the I/O nodes. This information is encapsulated in a `QIO_FileSystem` structure, which must be completed by the applications programmer.

```

typedef struct {
    int number_io_nodes;
    int type;
    int (*my_io_node)(const int node);
    int (*master_io_node)(void);
    int *io_node;
    char **node_path;
} QIO_FileSystem;

```

The `number_io_nodes` member specifies the number of I/O nodes. If it is the same as the `number_of_nodes` member of the layout structure, each node does its own I/O.

The `type` member is either `QIO_SINGLE_PATH` or `QIO_MULTI_PATH`. In single-path mode, all files are found in the same directory. In multi-path mode, a separate directory is specified for each I/O node.

The `my_io_node` function maps a node to its I/O node, based on the logical node number (rank). The `master_io_node` function returns the number of the master node.

The `io_node` table lists the numbers of the I/O nodes. If the number of I/O nodes is the same as the number of nodes, this table is not required, since each node does its own I/O. In that case the `my_io_node` function should be the identity map.

The `node_path` table is required only in multi-path mode. It lists the directories where the files for the I/O nodes are to be placed. The table has one entry for each I/O node. The entries must correlate with the entries in the `io_node` table.

Convert single file to partition file format

Prototype	<code>int QIO_single_to_part(const char filename[], QIO_FileSystem *fs, QIO_Layout *layout);</code>
Purpose	Convert an existing file from single to partition format.
Example	<code>QIO_single_to_part(filename, fs, mpp_layout);</code>

Convert partition file to single file format

Prototype	<code>int QIO_part_to_single(const char filename[], QIO_FileSystem *fs, QIO_Layout *layout);</code>
Purpose	Convert an existing file from single to partition format.
Example	<code>QIO_part_to_single(filename, fs, mpp_layout);</code>

As a matter of convenience, the file conversion application may be designed so that the code gets the lattice dimension and size from the file. The file should be opened by `QIO_open_read` with the layout `latdim` member set to zero. The lattice dimensions are then taken from the file and kept with the `QIO_Reader` structure. Two accessor utilities are provided for extracting the dimensions from the opaque structure, as described above.

4.6 String Handling with QIO

A few utilities are provided for manipulating the QIO string type `QIO.String` required by the API.

Creating an empty QIO String

Prototype	<code>QIO.String *QIO_string_create(void);</code>
Purpose	Creates an empty string.
Example	<code>fileinfo = QIO_string_create();</code>

Filling a QIO string from a null-terminated character array

Prototype	<code>QIO.String *QIO_string_set(QIO.String — *qs, const char *const string);</code>
Purpose	Inserts the null-terminated character array <code>string</code> into the string <code>qs</code> .
Example	<code>QIO_string *recinfo = QIO_string_create(); QIO_string_set(recinfo,string);</code>

Copying a QIO string

Prototype	<code>QIO.String *QIO_string_copy(QIO.String *dest, QIO.String *src);</code>
Purpose	Copies the string.
Example	<code>QIO_string_copy(newxml,oldxml);</code>

Resizing a string

Prototype	<code>QIO.String *QIO_string_realloc(QIO.String *dest, int length);</code>
Purpose	Change the length of the string with truncation if necessary.
Example	<code>QIO_string_realloc(xml,32);</code>

Accessing the string length

Prototype	<code>size_t QIO_string_bytes(const QIO.String *const xml);</code>
Purpose	Returns a pointer to the null-terminated character array in the string.
Example	<code>printf("%s\n", QIO_string_bytes(xml));</code>

Accessing the string character array

Prototype	<code>char *QIO_string_ptr(const QIO.String *const xml);</code>
Purpose	Returns the length of the string.
Example	<code>length = QIO_string_length(xml);</code>

Destroying a QIO string

Prototype	<code>void QIO_string_destroy(QIO_String *xml);</code>
Purpose	Frees storage.
Example	<code>QIO_string_destroy(xml);</code>

4.7 Compilation with QIO

There is a single top-level header file `qio.h` and a single library `libqio.a`. The QIO package is currently built in conjunction with the independent LIME package through `configure`, `make` and `make install`.